

EVSP DR Current Research

Updated 12 September 2026, 00:20 EDT. Results below are completed observations; overnight extensions are still running.

Current finding. With set covering and inherited columns, all six chains match the GIRO fleet at $k=3-6$. Chains 3, 5 and 6 also reach ten buses at $k=10$. Larger warm-chain cases remain unfinished.

Integer fleet with inherited columns

Each cell is the number of integer bus routes found. Green matches the target. Red requires extra buses. Grey means no completed result.

Target buses	Chain 1	Chain 2	Chain 3	Chain 4	Chain 5	Chain 6
2	2	2	2	2	3	2
3	3	3	3	3	3	3
4	4	4	4	4	4	4
5	5	5	5	5	5	5
6	6	6	6	6	6	6
7	—	7	7	7	7	7
8	—	8	8	8	8	8
9	—	—	9	9	9	9
10	—	—	10	—	10	10

k=11–15: no completed warm results yet. The overnight extension is queued through $k=15$ for all six chains.

$k=2$ starts from single-trip routes; later sizes reuse validated previous- k columns. Chain 3 $k=10$ also includes routes from a fresh solver solution. Displayed warm fleet optima are proved within their saved pools. Individual route replay passed; duplicate-trip removal and shared charger capacity are not established by this table.

Fresh starts on the same chains

Each size starts independently. Both tables use covering and 240-kWh batteries with 240-kW charging. Fresh fleet proof status varies by cell.

Target buses	Chain 1	Chain 2	Chain 3	Chain 4	Chain 5	Chain 6
5	5	5	5	5	6	5
8	9	9	9	9	9	8
10	11	11	11	11	11	11
15	18	17	18	19	16	19

Inherited columns improve the available integer pool: chain 3 at $k=8$ improves from a proved nine-bus pool optimum to eight; chain 5 at $k=5$ improves from a proved six-bus pool optimum to five.

What is limiting progress

Baseline warm runs: replaying every inherited route exhausted initialization budgets in chain 1 $k=7$, chain 2 $k=9$ and chain 4 $k=10$. The new treatment replays at most 512 sequences for 15 minutes, then starts ordinary CG. Compare it separately from full-pool inheritance.

Capacity pilot: pricing dominates. One $k=3$ pricing call took 7.15 hours; its LP took 0.006 seconds. CG stopped without a pricing certificate.

Case	Baseline	PARX 60 kW	Station capacity	Both
$k=1$, two duties tested	1 each	1 each	1 each	1 each
$k=2$, 23 trips	2	2	3	3
$k=3$, 35 trips	3	3	16	16

Sixteen is the result in a small generated pool, not a proved physical requirement.

Capacity-constrained schedules passed station sweeps. This pilot has no 65% return-SOC floor and is separate from the six baseline chains.

Running overnight

- 37 warm CG cases continue six chains to $k=15$.
- 40 independent CG cases split one 32-duty, 750-trip parent into four groups of eight, using ten different partitions.
- 10 recombination cases retain the component solution and try routes across groups.

Each CG case has a dependent one-hour, two-stage MIP on the default partition. At the check, 46 CG jobs were running. No completed result from this new campaign is shown above.

Methods and evidence

Each CG route costs 100,000 + electricity + 5 per charging start. MIP stage 1 minimizes buses. Stage 2 constrains buses $\leq$ the validated stage-1 incumbent and minimizes electricity plus start fees.

A pricing certificate establishes the LP result in the represented graph at its stated tolerance. An unfinished restricted LP objective is not a full-model lower bound. A MIP proof concerns its saved columns. Fleet target attainment, physical replay and shared-capacity validation are separate checks.

[Experiment register and exact source records](#) · [Overnight experiment settings](#) · [Historical document and figures](#)

Storage maintenance, 12 September: 156 cold files were archived losslessly, freeing 133.74 GB. Active research and V2G data were preserved. [Verification and restore instructions](#).